

克隆 ceph2 ceph3



```
[root@ceph1 ~]# cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain6
192.168.70.10 ceph1
192.168.70.20 ceph2
192.168.70.30 ceph3
192.168.70.40 ceph4
[root@ceph1 ~]#
```

```
[root@ceph1 ~]# cat /etc/sudoers.d/cephduser
cephduser ALL = (root) NOPASSWD:ALL
[root@ceph1 ~]#
```

systemctl restart chronyd

systemctl enable chronyd

chronyc sources -v

其中一台作为时钟源服务器 ceph1

```
# Use public servers from the pool.ntp.org project.
# Please consider joining the pool (http://www.pool.ntp.org/join.html).
#server 0.centos.pool.ntp.org iburst
#server 1.centos.pool.ntp.org iburst
#server 2.centos.pool.ntp.org iburst
#server 3.centos.pool.ntp.org iburst
server 192.168.70.10 iburst
allow 192.168.70.0/24
local stratum 10
```

```
[root@ceph1 ~]# chronyc sources -v
210 Number of sources = 1

... Source mode '^' = server, '=' = peer, '#' = local clock.
/.. Source state '*' = current synced, '+' = combined , '-' = not combined,
/ / '?' = unreachable, 'x' = time may be in error, '~' = time too variable.
||
|| Reachability register (octal) -. | - xxxx [ yyyy ] +/- zzzz
|| Log2(Polling interval) --. | | xxxx = adjusted offset,
|| | | yyyy = measured offset,
|| | | zzzz = estimated error.
|| | |
|| MS Name/IP address Stratum Poll Reach LastRx Last sample
=====
^* ceph1 10 6 7 0 -269ns[-2114ns] +/- 4919ns
[root@ceph1 ~]#
```

```
[root@ceph2 ~]# chronyc sources -v
210 Number of sources = 1

... Source mode '^' = server, '=' = peer, '#' = local clock.
/.. Source state '*' = current synced, '+' = combined , '-' = not combined,
/ / '?' = unreachable, 'x' = time may be in error, '~' = time too variable.
||
|| Reachability register (octal) -. | - xxxx [ yyyy ] +/- zzzz
|| Log2(Polling interval) --. | | xxxx = adjusted offset,
|| | | yyyy = measured offset,
|| | | zzzz = estimated error.
|| | |
|| MS Name/IP address Stratum Poll Reach LastRx Last sample
=====
^* ceph1 11 8 77 35 +34us[ +45us] +/- 23ms
[root@ceph2 ~]#
```

```
[root@ceph3 ~]# chronyc sources -v
210 Number of sources = 1

... Source mode '^' = server, '=' = peer, '#' = local clock.
/.. Source state '*' = current synced, '+' = combined , '-' = not combined,
| / '?' = unreachable, 'x' = time may be in error, '~' = time too variable.
||
|| Reachability register (octal) -. | xxxx [ yyyy ] +/- zzzz
|| Log2(Polling interval) --. | | xxxx = adjusted offset,
|| | yyyy = measured offset,
|| | zzzz = estimated error.
||
||
MS Name/IP address          Stratum Poll Reach LastRx Last sample
=====
^* ceph1                    11    7   377   100   +50us[ +43us] +/-  22ms
[root@ceph3 ~]#
```

```
[root@ceph4 ~]# chronyc sources -v
210 Number of sources = 1

... Source mode '^' = server, '=' = peer, '#' = local clock.
/.. Source state '*' = current synced, '+' = combined , '-' = not combined,
| / '?' = unreachable, 'x' = time may be in error, '~' = time too variable.
||
|| Reachability register (octal) -. | xxxx [ yyyy ] +/- zzzz
|| Log2(Polling interval) --. | | xxxx = adjusted offset,
|| | yyyy = measured offset,
|| | zzzz = estimated error.
||
||
MS Name/IP address          Stratum Poll Reach LastRx Last sample
=====
^* ceph1                    11    8    77    38   -109us[ -133us] +/-  23ms
[root@ceph4 ~]#
```

免密 ssh-keygen

ssh-copy-id cephuser@ceph1

ssh-copy-id cephuser@ceph2

ssh-copy-id cephuser@ceph3

```
[root@ceph1 ~]# ssh ceph2
Last login: Mon Jun 16 01:48:04 2025 from 192.168.70.1
[root@ceph2 ~]# exit
logout
Connection to ceph2 closed.
[root@ceph1 ~]# ceph 3
no valid command found; 10 closest matches:
config get <who> {<key>}
config dump
config set <who> <name> <value> {--force}
config rm <who> <name>
mgr count-metadata <property>
mgr versions
osd crush reweight <name> <float[0.0-]>
osd crush reweight-all
osd crush unlink <name> {<ancestor>}
osd crush rm <name> {<ancestor>}
Error EINVAL: invalid command
[root@ceph1 ~]# ceph3
-bash: ceph3: command not found
[root@ceph1 ~]# ssh ceph3
Last login: Mon Jun 16 01:48:15 2025 from 192.168.70.1
[root@ceph3 ~]# exit
logout
Connection to ceph3 closed.
[root@ceph1 ~]#
```

三节点一样生成 Ceph yum 源文件

```
[root@ceph1 ~]# cd /etc/yum.repos.d/
[root@ceph1 yum.repos.d]# cat /etc/yum.repos.d/ceph.repo
[Ceph]
name=Ceph packages for
baseurl=https://mirrors.tuna.tsinghua.edu.cn/ceph/rpm-nautilus/el7/x86_64/

enabled=1
gpgcheck=1
type=rpm-md
gpgkey=https://download.ceph.com/keys/release.asc
priority=1

[Ceph-noarch]
name=Ceph noarch packages
baseurl=https://mirrors.tuna.tsinghua.edu.cn/ceph/rpm-nautilus/el7/noarch
enabled=1
gpgcheck=1
type=rpm-md
gpgkey=https://download.ceph.com/keys/release.asc
priority=1

[ceph-source]
name=Ceph source packages
baseurl=https://mirrors.tuna.tsinghua.edu.cn/ceph/rpm-nautilus/el7/SRPMS
enabled=1
gpgcheck=1
type=rpm-md
gpgkey=https://download.ceph.com/keys/release.asc
[root@ceph1 yum.repos.d]#
```

ssh://192.168.70.10:22 SSH2 xterm 118x55 55,27 4 会话 CAP NUM

```
[root@ceph1 ceph_cluster]# ceph-deploy new ceph1 ceph2
Traceback (most recent call last):
  File "/usr/bin/ceph-deploy", line 5, in <module>
    from pkg_resources import load_entry_point
ImportError: No module named pkg_resources
[root@ceph1 ceph_cluster]# python get-pip.py
DEPRECATION: Python 2.7 reached the end of its life on January 1st, 2020. Please upgrade your Python as Python 2
.7 is no longer maintained. pip 21.0 will drop support for Python 2.7 in January 2021. More details about Python
2 support in pip can be found at https://pip.pypa.io/en/latest/development/release-process/#python-2-support pi
p 21.0 will remove support for this functionality.
Collecting pip<21.0
  Downloading pip-20.3.4-py2.py3-none-any.whl (1.5 MB)
    | 1.5 MB 21 kB/s
Collecting setuptools<45
  Downloading setuptools-44.1.1-py2.py3-none-any.whl (583 kB)
    | 583 kB 11 kB/s
Collecting wheel
```

```
[root@ceph1 ~]# yum list | grep ceph-deploy
Repodata is over 2 weeks old. Install yum-cron? Or run: yum makecache fast
ceph-deploy.noarch                2.0.1-0                @Ceph-noarch
[root@ceph1 ~]#
```

ssh://192.168.70.10:22 SSH2 xterm 126x31 31,17 4 会话 CAP NUM

ceph-deploy new ceph1 ceph2 ceph3

```
[root@ceph1 ceph_cluster]# ceph-deploy new ceph1 ceph2 ceph3
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy new ceph1 ceph2 ceph3
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username : None
```

```
[root@ceph1 ~]# ceph -s
cluster:
  id: 160331b1-12f9-4a84-bfde-cb6f2ce05282
  health: HEALTH_WARN
    Reduced data availability: 161 pgs inactive, 139 pgs peering
    208 pgs not deep-scrubbed in time
    208 pgs not scrubbed in time
    8 slow ops, oldest one blocked for 1549 sec, osd.0 has slow ops
    mons are allowing insecure global_id reclaim

services:
  mon: 3 daemons, quorum ceph1,ceph2,ceph3 (age 8m)
  mgr: ceph2(active, since 26m), standbys: ceph3, ceph1
  osd: 3 osds: 3 up (since 25m), 3 in (since 4w)

data:
  pools: 7 pools, 208 pgs
  objects: 245 objects, 37 MiB
  usage: 3.1 GiB used, 57 GiB / 60 GiB avail
  pgs: 77.404% pgs not active
    139 peering
    47 active+clean
    22 activating
```

vim ceph.conf 三台

```
[root@ceph1 ceph_cluster]# cat ceph.conf
[global]
fsid = 160331b1-12f9-4a84-bfde-cb6f2ce05282
mon_initial_members = ceph1, ceph2, ceph3
mon_host = 192.168.70.10,192.168.70.20,192.168.70.30
auth_cluster_required = cephx
auth_service_required = cephx
auth_client_required = cephx
public_network = 192.168.70.0/24
cluster_network = 172.18.1.0/24
```

后端网络 172.18.1.10 / 11 / 12

```
3: ens34: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen 1000
    link/ether 00:0c:29:95:e5:51 brd ff:ff:ff:ff:ff:ff
    inet 172.18.1.10/24 brd 172.18.1.255 scope global noprefixroute ens34
```

推送

ceph-deploy --overwrite-conf config push ceph1 ceph2

sudo systemctl restart ceph-\*.target

```
[root@ceph1 ceph_cluster]# sudo systemctl restart ceph-*.target
[root@ceph1 ceph_cluster]# yum list | grep ceph.x86_64
calceph.x86_64                3.5.3-2.el7                                epel
ceph.x86_64                   2:14.2.22-0.el7                            Ceph
collectd-ceph.x86_64          5.8.1-2.el7                                epel
[root@ceph1 ceph_cluster]# sudo yum install -y ceph
Loaded plugins: fastestmirror
Loading mirror speeds from cached hostfile
```

初始化 mon, 生成检测集群所需的密钥

```
[root@ceph1 ceph_cluster]# ceph-deploy mon create-initial
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy mon create-initial
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                 : False
[ceph_deploy.cli][INFO ] overwrite_conf          : False
[ceph_deploy.cli][INFO ] subcommand              : create-initial
[ceph_deploy.cli][INFO ] quiet                   : False
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7f46e462be18>
```

ceph-deploy mon create-initial

```
[root@ceph1 ceph_cluster]# ceph-deploy mon create-initial
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy mon create-initial
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                 : False
[ceph_deploy.cli][INFO ] overwrite_conf          : False
```

```
[root@ceph1 ceph_cluster]# ls
ceph.bootstrap-mds.keyring  ceph.bootstrap-osd.keyring  ceph.client.admin.keyring  ceph-deploy-ceph.log  y?
ceph.bootstrap-mgr.keyring  ceph.bootstrap-rgw.keyring  ceph.conf                  ceph.mon.keyring      y?.pub
[root@ceph1 ceph_cluster]#
```

ssh://192.168.70.10:22

sudo systemctl status ceph-mon@ceph1

```
[root@ceph1 ceph_cluster]# sudo systemctl status ceph-mon@ceph1
● ceph-mon@ceph1.service - Ceph cluster monitor daemon
   Loaded: loaded (/usr/lib/systemd/system/ceph-mon@ceph1.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2025-06-16 02:49:09 EDT; 38s ago
     Main PID: 2828 (ceph-mon)
    CGroup: /system.slice/system-ceph\x2dmon.slice/ceph-mon@ceph1.service
            └─2828 /usr/bin/ceph-mon -f --cluster ceph --id ceph1 --setuser ceph --setgroup ceph

Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:191 Unknown lvalue 'LockPersonality' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:201 Unknown lvalue 'MemoryDenyWriteExecute' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:241 Unknown lvalue 'ProtectControlGroups' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:261 Unknown lvalue 'ProtectKernelModules' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:271 Unknown lvalue 'ProtectKernelTunables' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:191 Unknown lvalue 'LockPersonality' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:201 Unknown lvalue 'MemoryDenyWriteExecute' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:241 Unknown lvalue 'ProtectControlGroups' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:261 Unknown lvalue 'ProtectKernelModules' in section 'Service'
Jun 16 02:49:23 ceph1 systemd[1]: /usr/lib/systemd/system/ceph-mon@ceph1.service:271 Unknown lvalue 'ProtectKernelTunables' in section 'Service'
[root@ceph1 ceph_cluster]#
```

ssh://192.168.70.10:22



sudo systemctl status ceph-mon@ceph2

```
[root@ceph2 yum.repos.d]# sudo systemctl status ceph-mon@ceph2
● ceph-mon@ceph2.service - Ceph cluster monitor daemon
   Loaded: loaded (/usr/lib/systemd/system/ceph-mon@.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2025-06-16 01:45:24 EDT; 1h 5min ago
     Main PID: 1014 (ceph-mon)
    CGroup: /system.slice/system-ceph\x2dmon.slice/ceph-mon@ceph2.service
            └─1014 /usr/bin/ceph-mon -f --cluster ceph --id ceph2 --setuser ceph --setgroup ceph

Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:19] Unknown lvalue 'LockPersonal...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:20] Unknown lvalue 'MemoryDenyWr...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:24] Unknown lvalue 'ProtectContr...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:26] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:27] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:19] Unknown lvalue 'LockPersonal...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:20] Unknown lvalue 'MemoryDenyWr...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:24] Unknown lvalue 'ProtectContr...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:26] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:49:26 ceph2 systemd[1]: [/usr/lib/systemd/system/ceph-mon@.service:27] Unknown lvalue 'ProtectKerne...rvice'
Hint: Some lines were ellipsized, use -l to show in full.
[root@ceph2 yum.repos.d]#
```

## 认证

```
[root@ceph1 ceph_cluster]# ceph-deploy admin ceph1 ceph2 ceph3
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy admin ceph1 ceph2 ceph3
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username           : None
[ceph_deploy.cli][INFO ] verbose           : False
[ceph_deploy.cli][INFO ] overwrite_conf    : False
[ceph_deploy.cli][INFO ] quiet             : False
[ceph_deploy.cli][INFO ] cd_conf           : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7ff775b6a248>
>
[ceph_deploy.cli][INFO ] cluster           : ceph
```

```
[root@ceph1 ceph_cluster]# ls /etc/ceph
ceph.client.admin.keyring  ceph.conf  rbdmap  tmpEmCCwQ
[root@ceph1 ceph_cluster]#
```

## Mgr

```
[root@ceph1 ceph_cluster]# ceph-deploy mgr create ceph1 ceph2
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy mgr create ceph1 ceph2
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username           : None
[ceph_deploy.cli][INFO ] verbose           : False
[ceph_deploy.cli][INFO ] mgr                : [('ceph1', 'ceph1'), ('ceph2', 'ceph2')]
[ceph_deploy.cli][INFO ] overwrite_conf     : False
[ceph_deploy.cli][INFO ] subcommand         : create
[ceph_deploy.cli][INFO ] quiet              : False
[ceph_deploy.cli][INFO ] cd_conf            : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7f48d4f5095>
0>
[ceph_deploy.cli][INFO ] cluster           : ceph

[ceph2][INFO ] Running command: systemctl enable ceph.target
[root@ceph1 ceph_cluster]# sudo ss -antlup | grep mgr
tcp    LISTEN  0      128      192.168.70.10:6804      *:*                users:({"ceph-mgr",pid=1019,fd=26})
tcp    LISTEN  0      128      192.168.70.10:6805      *:*                users:({"ceph-mgr",pid=1019,fd=29})
[root@ceph1 ceph_cluster]#
```

```

[10-257]
[root@ceph1 ceph_cluster]# ceph -s
cluster:
  id:         160331b1-12f9-4a84-bfde-cb6f2ce05282
  health:     HEALTH_WARN
              application not enabled on 1 pool(s)
              mons are allowing insecure global_id reclaim

services:
  mon: 3 daemons, quorum ceph1,ceph2,ceph3 (age 12m)
  mgr: ceph1(active, since 11m), standbys: ceph2, ceph3
  osd: 6 osds: 6 up (since 7m), 6 in (since 7m)
  rgw: 1 daemon active (ceph3)

task status:

data:
  pools:   8 pools, 224 pgs
  objects: 249 objects, 38 MiB
  usage:   6.2 GiB used, 84 GiB / 90 GiB avail
  pgs:     224 active+clean

```

ceph-deploy osd create --data /dev/sdb ceph1

ceph-deploy osd create --data /dev/sdb ceph2

ceph-deploy osd create --data /dev/sdb ceph3

检查 osd 服务状态

- 能查询到正确的 osd 服务数量，并且都处于 up 和 in 状态

```

[root@ceph1 ceph_cluster]# ceph -s
cluster:
  id:         160331b1-12f9-4a84-bfde-cb6f2ce05282
  health:     HEALTH_WARN
              Degraded data redundancy: 245/735 objects degraded (33.333%), 62 pgs degraded, 208 pgs undersized
              208 pgs not deep-scrubbed in time
              208 pgs not scrubbed in time
              mons are allowing insecure global_id reclaim

services:
  mon: 3 daemons, quorum ceph1,ceph2,ceph3 (age 10m)
  mgr: ceph2(active, since 74m), standbys: ceph3, ceph1
  osd: 3 osds: 2 up (since 10m), 2 in (since 16m)
  rgw: 1 daemon active (ceph3)

```

```

[root@ceph1 ceph_cluster]# ceph df
RAW STORAGE:
  CLASS  SIZE      AVAIL      USED      RAW USED    %RAW USED
  hdd    40 GiB    38 GiB    85 MiB    2.1 GiB      5.21
  TOTAL  40 GiB    38 GiB    85 MiB    2.1 GiB      5.21

POOLS:
  POOL          ID  PGS  STORED  OBJECTS  USED      %USED  MAX AVAIL
  rpdpool       2   16    49 MiB    19      100 MiB    0.18    27 GiB
  .rgw.root     3   32    1.8 KiB    4       768 KiB    0       27 GiB
  default.rgw.control  4   32    0 B       8        0 B      0       27 GiB
  default.rgw.meta     5   32    1.2 KiB    5       768 KiB    0       27 GiB
  default.rgw.log       6   32    0 B       207      0 B      0       27 GiB
  default.rgw.buckets.index  7   32    221 B     1       221 B    0       27 GiB
  default.rgw.buckets.data  8   32    24 B     1       192 KiB  0       27 GiB

```

ceph health detail

```

120 more pgs...
SLOW_OPS 8 slow ops, oldest one blocked for 1614 sec, osd.0 has slow ops
AUTH_INSECURE_GLOBAL_ID_RECLAIM_ALLOWED mons are allowing insecure global_id reclaim
  mon.ceph1 has auth_allow_insecure_global_id_reclaim set to true
  mon.ceph2 has auth_allow_insecure_global_id_reclaim set to true
  mon.ceph3 has auth_allow_insecure_global_id_reclaim set to true
[root@ceph1 ~]#

```

```
[root@ceph1 ceph_cluster]# ceph osd tree
ID CLASS WEIGHT  TYPE NAME        STATUS REWEIGHT PRI-AFF
-1              0.05846  root default
-3              0.01949   host ceph1
 0 hdd 0.01949     osd.0        up  1.00000 1.00000
-5              0.01949   host ceph2
 1 hdd 0.01949     osd.1        up  1.00000 1.00000
-7              0.01949   host ceph3
 2 hdd 0.01949     osd.2        up  1.00000 1.00000
[root@ceph1 ceph_cluster]#
```

ssh://192.168.70.10:22 SSH2 xterm 118x55 55,28 4 会话 CAP NUM

```
[root@ceph1 ~]# ceph mon stat
el: 3 mons at {ceph1=[v2:192.168.70.10:3300/0,v1:192.168.70.10:6789/0],ceph2=[v2:192.168.70.20:3300/0,v1:192.168.70.20:6789/0],ceph3=[v2:192.168.70.30:3300/0,v1:192.168.70.30:6789/0]}, election epoch 34, leader 0 ceph1, quorum 0,1,2 ceph1,ceph2,ceph3
[root@ceph1 ~]# ceph mon stat
el: 3 mons at {ceph1=[v2:192.168.70.10:3300/0,v1:192.168.70.10:6789/0],ceph2=[v2:192.168.70.20:3300/0,v1:192.168.70.20:6789/0],ceph3=[v2:192.168.70.30:3300/0,v1:192.168.70.30:6789/0]}, election epoch 34, leader 0 ceph1, quorum 0,1,2 ceph1,ceph2,ceph3
[root@ceph1 ~]# ceph osd stat
3 osds: 3 up (since 27m), 3 in (since 4w); epoch: e63
[root@ceph1 ~]#
```

ssh://192.168.70.10:22 SSH2 xterm 126x31 31,17 4 会话 CAP NUM

### ceph-deploy disk list

```
[root@ceph1 ceph_cluster]# ceph-deploy disk list ceph1
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.ceph_deploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy disk list ceph1
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] debug                  : False
[ceph_deploy.cli][INFO ] overwrite_conf         : False
[ceph_deploy.cli][INFO ] subcommand              : list
[ceph_deploy.cli][INFO ] quiet                  : False
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph1>
0>
[ceph_deploy.cli][INFO ] cluster                 : ceph1
[ceph_deploy.cli][INFO ] host                     : ['ceph1']
[ceph_deploy.cli][INFO ] func                     : <func>
[ceph_deploy.cli][INFO ] ceph_conf                : None
[ceph_deploy.cli][INFO ] default_release         : Fedora
[ceph1][DEBUG ] connected to host: ceph1
[ceph1][DEBUG ] detect platform information from remote host
[ceph1][DEBUG ] detect machine type
[ceph1][DEBUG ] find the location of an executable
[ceph1][INFO ] Running command: fdisk -l
[ceph1][INFO ] Disk /dev/sda: 21.5 GB, 21474836480 bytes, 419430400 sectors
[ceph1][INFO ] Disk /dev/sdb: 21.5 GB, 21474836480 bytes, 419430400 sectors
[ceph1][INFO ] Disk /dev/mapper/centos-root: 18.2 GB, 182494720 sectors
[ceph1][INFO ] Disk /dev/mapper/centos-swap: 2147 MB, 2147483648 bytes, 419430400 sectors
[ceph1][INFO ] Disk /dev/mapper/ceph--b563822c--dc3d--444d--8f328c2f6f4: 21.5 GB, 21470642176 bytes, 41934848 sectors
```



```

[root@ceph1 ceph_cluster]# ceph-deploy disk list ceph2
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.ceph/ceph-deploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] debug                  : False
[ceph_deploy.cli][INFO ] overwrite_conf         : False
[ceph_deploy.cli][INFO ] subcommand             : list
[ceph_deploy.cli][INFO ] quiet                  : False
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph>
0>
[ceph_deploy.cli][INFO ] cluster                : ceph
[ceph_deploy.cli][INFO ] host                   : ['ceph2']
[ceph_deploy.cli][INFO ] func                   : <func>
[ceph_deploy.cli][INFO ] ceph_conf              : None
[ceph_deploy.cli][INFO ] default_release        : False
[ceph2][DEBUG ] connected to host: ceph2
[ceph2][DEBUG ] detect platform information from remote host
[ceph2][DEBUG ] detect machine type
[ceph2][DEBUG ] find the location of an executable
[ceph2][INFO ] Running command: fdisk -l
[ceph2][INFO ] Disk /dev/sdb: 21.5 GB, 21474836480 bytes, 4194304 sectors
[ceph2][INFO ] Disk /dev/sda: 21.5 GB, 21474836480 bytes, 4194304 sectors
[ceph2][INFO ] Disk /dev/mapper/centos-root: 18.2 GB, 1824941696 bytes, 3568000 sectors
[ceph2][INFO ] Disk /dev/mapper/centos-swap: 2147 MB, 2147483648 bytes, 4194304 sectors
[ceph2][INFO ] Disk /dev/mapper/ceph--26bee39d--7da1--4bbe--884
--7a60427489be: 21.5 GB, 21470642176 bytes, 41934848 sectors

```

```

[root@ceph1 ceph_cluster]# ceph-deploy disk list ceph3
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.ceph/ceph-deploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] debug                  : False
[ceph_deploy.cli][INFO ] overwrite_conf         : False
[ceph_deploy.cli][INFO ] subcommand             : list
[ceph_deploy.cli][INFO ] quiet                  : False
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph>
0>
[ceph_deploy.cli][INFO ] cluster                : ceph
[ceph_deploy.cli][INFO ] host                   : ['ceph3']
[ceph_deploy.cli][INFO ] func                   : <func>
[ceph_deploy.cli][INFO ] ceph_conf              : None
[ceph_deploy.cli][INFO ] default_release        : False
[ceph3][DEBUG ] connected to host: ceph3
[ceph3][DEBUG ] detect platform information from remote host
[ceph3][DEBUG ] detect machine type
[ceph3][DEBUG ] find the location of an executable
[ceph3][INFO ] Running command: fdisk -l
[ceph3][INFO ] Disk /dev/sda: 21.5 GB, 21474836480 bytes, 4194304 sectors
[ceph3][INFO ] Disk /dev/sdb: 21.5 GB, 21474836480 bytes, 4194304 sectors
[ceph3][INFO ] Disk /dev/mapper/centos-root: 18.2 GB, 1824941696 bytes, 3568000 sectors
[ceph3][INFO ] Disk /dev/mapper/centos-swap: 2147 MB, 2147483648 bytes, 4194304 sectors
[ceph3][INFO ] Disk /dev/mapper/ceph--e333a0f7--640b--4db5-
--7f6b88a26fel: 21.5 GB, 21470642176 bytes, 41934848 sectors

```

再加三块三台都三块



硬盘 (SCSI)

20 GB

硬盘 2 (SCSI)

20 GB

硬盘 3 (SCSI)

10 GB

此虚拟机的内存(M):

4096 MB

64 GB

```

[root@ceph1 ~]# lsblk
NAME                                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
sda                                  8:0    0   20G  0 disk 
├─sda1                              8:1    0    1G  0 part /boot
├─sda2                              8:2    0   19G  0 part 
│   └─centos-root                    253:0    0   17G  0 lvm  /
│       └─centos-swap                253:1    0    2G  0 lvm  [SWAP]
└─sdb                                8:16   0   20G  0 disk 
    └─ceph--b563822c--dc3d--444d--846f--9c2fff48f019-osd--block--c09db35c--a0bc--4744--a41e--f8f328c2f6f4 253:2    0   20G  0 lvm 
        └─sdc                        8:32    0   10G  0 disk 
            └─sr0                     11:0    1   973M  0 rom

```

Sdc

ceph-deploy disk zap ceph1 /dev/sdc

```

[root@ceph1 ~]# ceph-deploy disk zap ceph1 /dev/sdc
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy disk zap ceph1 /dev/sdc
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] username                : None
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] debug                  : False
[ceph_deploy.cli][INFO ] overwrite_conf         : False
[ceph_deploy.cli][INFO ] subcommand              : zap

```

ceph-deploy osd create --data /dev/sdc ceph1

ceph-deploy osd create --data /dev/sdc ceph2

ceph-deploy osd create --data /dev/sdc ceph3

```

[root@ceph1 ~]# cd ceph_cluster/
[root@ceph1 ceph_cluster]# ceph-deploy osd create --data /dev/sdc ceph1
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy osd create --data /dev/sdc ceph1
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] bluestore              : None
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7fc59ac9c68
0>
[ceph_deploy.cli][INFO ] cluster                 : ceph
[ceph_deploy.cli][INFO ] fs_type                 : xfs

```

```

[root@ceph1 ceph_cluster]# ceph-deploy osd create --data /dev/sdc ceph2
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy osd create --data /dev/sdc ceph2
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] bluestore              : None
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7ff8e16ea68
0>
[ceph_deploy.cli][INFO ] cluster                 : ceph
[ceph_deploy.cli][INFO ] fs_type                 : xfs

```

```

[root@ceph1 ceph_cluster]# ceph-deploy osd create --data /dev/sdc ceph3
[ceph_deploy.conf][DEBUG ] found configuration file at: /root/.cephdeploy.conf
[ceph_deploy.cli][INFO ] Invoked (2.0.1): /usr/bin/ceph-deploy osd create --data /dev/sdc ceph3
[ceph_deploy.cli][INFO ] ceph-deploy options:
[ceph_deploy.cli][INFO ] verbose                : False
[ceph_deploy.cli][INFO ] bluestore              : None
[ceph_deploy.cli][INFO ] cd_conf                 : <ceph_deploy.conf.cephdeploy.Conf instance at 0x7f6e61cdf68
0>
[ceph_deploy.cli][INFO ] cluster                 : ceph
[ceph_deploy.cli][INFO ] fs_type                 : xfs

```

查看

```

services:
mon: 3 daemons, quorum ceph1,ceph2,ceph3 (age 4m)
mgr: ceph1(active, since 4m), standbys: ceph2, ceph3
osd: 6 osds: 6 up (since 11s), 6 in (since 11s); 2 remapped pgs
rgw: 1 daemon active (ceph3)

```

Pool 列出

```
[root@ceph1 ceph_cluster]# ceph osd pool ls
rpdpool
.rgw.root
default.rgw.control
default.rgw.meta
default.rgw.log
default.rgw.buckets.index
default.rgw.buckets.data
[root@ceph1 ceph_cluster]#
```

sudo yum install ceph-mgr -y (ceph3 执行)

为集群添加 mgr 节点(主节点执行)

ceph-deploy mgr create ceph3

查询运行状态(节点 3 执行)

```
[root@ceph3 yum.repos.d]# sudo systemctl status ceph-mgr@ceph3.service
● ceph-mgr@ceph3.service - Ceph cluster manager daemon
   Loaded: loaded (/usr/lib/systemd/system/ceph-mgr@.service; enabled; vendor preset: disabled)
   Active: active (running) since Mon 2025-06-16 01:45:27 EDT; 1h 17min ago
     Main PID: 1015 (ceph-mgr)
        CGroup: /system.slice/system-ceph-x2dmgr.slice/ceph-mgr@ceph3.service
                └─1015 /usr/bin/ceph-mgr -f --cluster ceph --id ceph3 --setuser ceph --setgroup ceph

Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:15] Unknown lvalue 'LockPersonal...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:18] Unknown lvalue 'MemoryDenyWr...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:21] Unknown lvalue 'ProtectContr...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:23] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:24] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:15] Unknown lvalue 'LockPersonal...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:18] Unknown lvalue 'MemoryDenyWr...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:21] Unknown lvalue 'ProtectContr...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:23] Unknown lvalue 'ProtectKerne...rvice'
Jun 16 02:59:07 ceph3 systemd[1]: [/usr/lib/systemd/system/ceph-mgr@.service:24] Unknown lvalue 'ProtectKerne...rvice'
Hint: Some lines were ellipsized, use -l to show in full.
[root@ceph3 yum.repos.d]#
```

RBD

主节点

```
[root@ceph1 ~]# scp -p /etc/ceph/ceph.conf root@192.168.70.40:/etc/ceph/
root@192.168.70.40's password:
ceph.conf                                100% 303    78.3KB/s   00:00
[root@ceph1 ~]# scp -p /etc/ceph/ceph.client.admin.keyring root@192.168.70.40:/etc/ceph
root@192.168.70.40's password:
ceph.client.admin.keyring                100% 151    141.4KB/s   00:00
[root@ceph1 ~]#
```

Ceph4

```
[root@ceph4 ~]# yum install -y ceph-common.x86_64
Loaded plugins: fastestmirror
Determining fastest mirrors
Could not get metalink https://mirrors.fedoraproject.org/metalink?repo=epel-7&arch=x86_64 error was
12: Timeout on https://mirrors.fedoraproject.org/metalink?repo=epel-7&arch=x86_64: (28, 'Operation timed out after 300
01 milliseconds with 0 out of 0 bytes received')
 * base: mirrors.aliyun.com
 * epel: d2lzk17pfhq30w.cloudfront.net
 * extras: mirrors.aliyun.com
```

```

[root@ceph4 ~]# rbd map rbdpool/sqlldb
/dev/rbd0
[root@ceph4 ~]# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
sda           8:0    0   20G  0 disk
├─sda1        8:1    0    1G  0 part /boot
├─sda2        8:2    0   19G  0 part
│ └─centos-root 253:0    0   17G  0 lvm  /
│   └─centos-swap 253:1    0    2G  0 lvm  [SWAP]
sdb           8:16    0   20G  0 disk
sr0          11:0    1  973M  0 rom
rbd0         252:0    0    1G  0 disk
[root@ceph4 ~]# mkfs.ext4 /dev/rbd0
mkfs2fs 1.42.9 (28-Dec-2013)
Discarding device blocks: done
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=1024 blocks, Stripe width=1024 blocks
65536 inodes, 262144 blocks
13107 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=268435456
8 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376

Allocating group tables: done
Writing inode tables: done
Creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done

```

在 ceph4 创建挂载

```

[root@ceph4 ~]# mkdir /mnt/rbd
[root@ceph4 ~]# mount /dev/rbd0 /mnt/rbd
[root@ceph4 ~]# df -h
Filesystem                Size      Used Avail Use% Mounted on
devtmpfs                  1.9G         0  1.9G   0% /dev
tmpfs                     1.9G         0  1.9G   0% /dev/shm
tmpfs                     1.9G      12M  1.9G   1% /run
tmpfs                     1.9G         0  1.9G   0% /sys/fs/cgroup
/dev/mapper/centos-root    17G       1.9G   16G  11% /
/dev/sda1                 1014M     151M   864M  15% /boot
tmpfs                     378M         0   378M   0% /run/user/0
/dev/rbd0                 976M     2.6M   907M   1% /mnt/rbd
[root@ceph4 ~]#

```

写入赵文字 +230303060 +云计算 2 班

```

[root@ceph4 ~]# echo "赵文字 +230303060 +云计算 2班" > /mnt/rbd/student_info.txt
[root@ceph4 ~]# cat /mnt/rbd/student_info.txt
赵文字 +230303060 +云计算 2班
[root@ceph4 ~]# df -h
Filesystem                Size      Used Avail Use% Mounted on
devtmpfs                  1.9G         0  1.9G   0% /dev
tmpfs                     1.9G         0  1.9G   0% /dev/shm
tmpfs                     1.9G      12M  1.9G   1% /run
tmpfs                     1.9G         0  1.9G   0% /sys/fs/cgroup
/dev/mapper/centos-root    17G       1.9G   16G  11% /
/dev/sda1                 1014M     151M   864M  15% /boot
tmpfs                     378M         0   378M   0% /run/user/0
/dev/rbd0                 976M     2.6M   907M   1% /mnt/rbd
[root@ceph4 ~]#

```